lab4

Hugo Morvan, Daniele Bozzoli

2023-11-28

Computational Statistics - Fall 23

Computer Lab 4

Hugo Morvan, Daniele Bozzoli

Question 1: Computations with Metropolis-Hastings

Consider a random variabel X with the following probability density function: $f(x) = x5e - x, x > 0$. The distribution is known up to some constant of proportionality. If you are interested (NOT part of the Lab) this constant can be found by applying integration by parts multiple times and equals 120.

a.

Use the Metropolis–Hastings algorithm to generate 10000 samples from this distribution by using a log–normal LN (Xt , 1) proposal distribution; take some starting point. Plot the chain you obtained with iterations on the horizontal axis. What can you guess about the convergence of the chain? If there is a burn–in period, what can be the size of this period? What is the acceptance rate? Plot a histogram of the sample.

b.

Perform Part a by using the chi–square distribution as a proposal distribution, where is the floor function, meaning the integer part of x for positive x

c.

Suggest another proposal distribution (can be a log normal or chi–square distribution with other parameters or another distribution) with the potential to generate a good sample. Perform part a with this distribution.

d.

Compare the results of Parts a, b, and c and make conclusions.

e.

Estimate

$$E(X) = \int_0^\infty x f(x) dx$$

using the samples from Parts a, b, and c.

f.

The distribution generated is in fact a gamma distribution. Look in the literature and define the actual value of the integral. Compare it with the one you obtained.

Question 2: Gibbs sampling

Let X = (X1 , X2) be a bivariate distribution with density $f(x_1, x_2) = 1\{x_1^2 + wx_1x_2 + x_2^2 < 1\}$ for some specific w with |w| < 2. X has a uniform distribution on some two-dimensional region. We consider here the case w = 1.999 (in Lecture 4, the case w = 1.8 was shown).

a.

Draw the boundaries of the region where X has a uniform distribution. You can use the code provided on the course homepage and adjust it.

```
################################
### Boundary ellipse
### for Gibbs sampling example
### from Lecture 4
### Fall 2023, by Frank Miller
################################
w  <- 1.999
xv <- seq(-1, 1, by=0.01) * 1/sqrt(1-w^2/4)  # a range of x1-values, where the term below the root is n
plot(xv, xv, type="n", xlab=expression(x[1]), ylab=expression(x[2]), las=1)
# ellipse
lines(xv, -(w/2)*xv-sqrt(1-(1-w^2/4)*xv^2), lwd=2, col=8)
lines(xv, -(w/2)*xv+sqrt(1-(1-w^2/4)*xv^2), lwd=2, col=8)
```

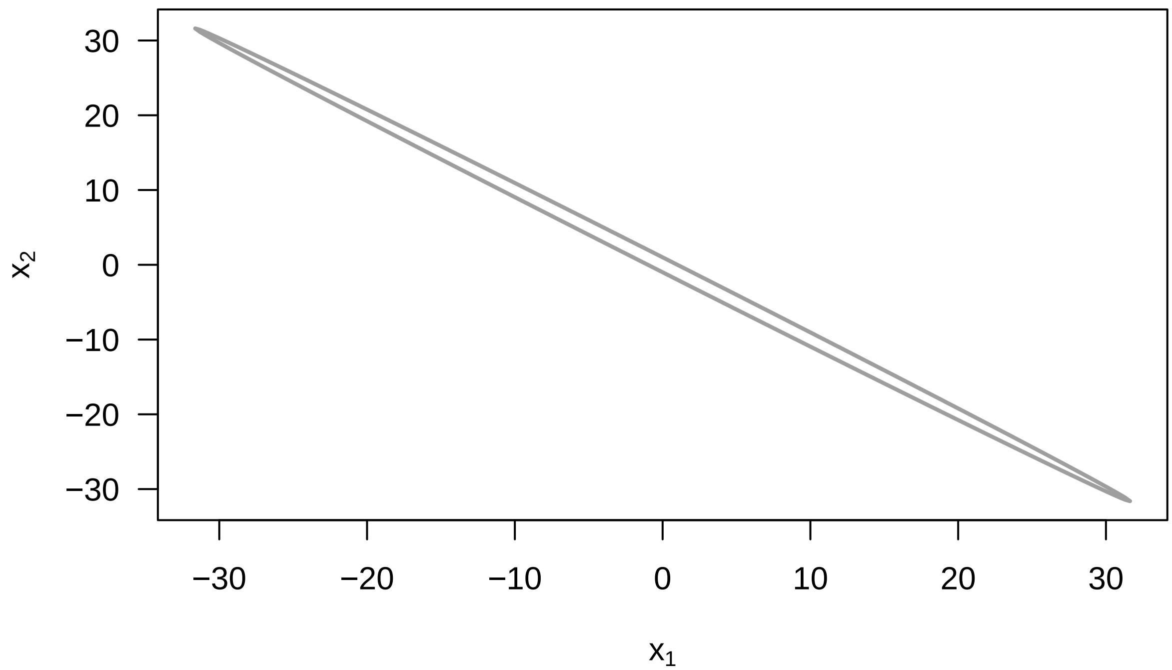

b.

What is the conditional distribution of X1 given X2 and that of X2 given X1 ?

The conditional distribution $f(x_1|x_2)$ is a uniform distribution on the interval $(-0.9995x_1 - \sqrt{1-0.00099975x_1}, -0.9995x_1 + \sqrt{1-0.00099975x_1})$

The conditional distribution $f(x_2|x_1)$ is a uniform distribution on the interval $(-0.9995x_2 - \sqrt{1-0.00099975x_2}, -0.9995x_2 + \sqrt{1-0.00099975x_2})$

c.

Write your own code for Gibbs sampling the distribution. Run it to generate n = 1000 random vectors and plot them into the picture from Part a.

```
#Gibbs sampling
f1_cond <- function(x2){
  return(runif(1, min = -0.9995*x2 - sqrt(1-0.00099975*x2), max = -0.9995*x2 + sqrt(1-0.00099975*x2)))
}
f2_cond <- function(x1){
  return(runif(1, min = -0.9995*x1 - sqrt(1-0.00099975*x1), max = -0.9995*x1 + sqrt(1-0.00099975*x1)))
}

set.seed(12345)
n <- 1000
```

```r
x1 <- 0
x2 <- 0
X1 <- c()
X2 <- c()
for(i in 1:n){
  x1 <- f1_cond(x2)
  x2 <- f2_cond(x1)
  X1 <- c(X1, x1)
  X2 <- c(X2, x2)
}

w  <- 1.999
xv <- seq(-1, 1, by=0.01) * 1/sqrt(1-w^2/4)  # a range of x1-values, where the term below the root is n
plot(xv, xv, type="n", xlab=expression(x[1]), ylab=expression(x[2]), las=1)
# ellipse
lines(xv, -(w/2)*xv-sqrt(1-(1-w^2/4)*xv^2), lwd=2, col=8)
lines(xv, -(w/2)*xv+sqrt(1-(1-w^2/4)*xv^2), lwd=2, col=8)
#Add the points X1 and X2 to the graph
points(X1, X2, col = "red", pch = 20, cex = 0.5)
```

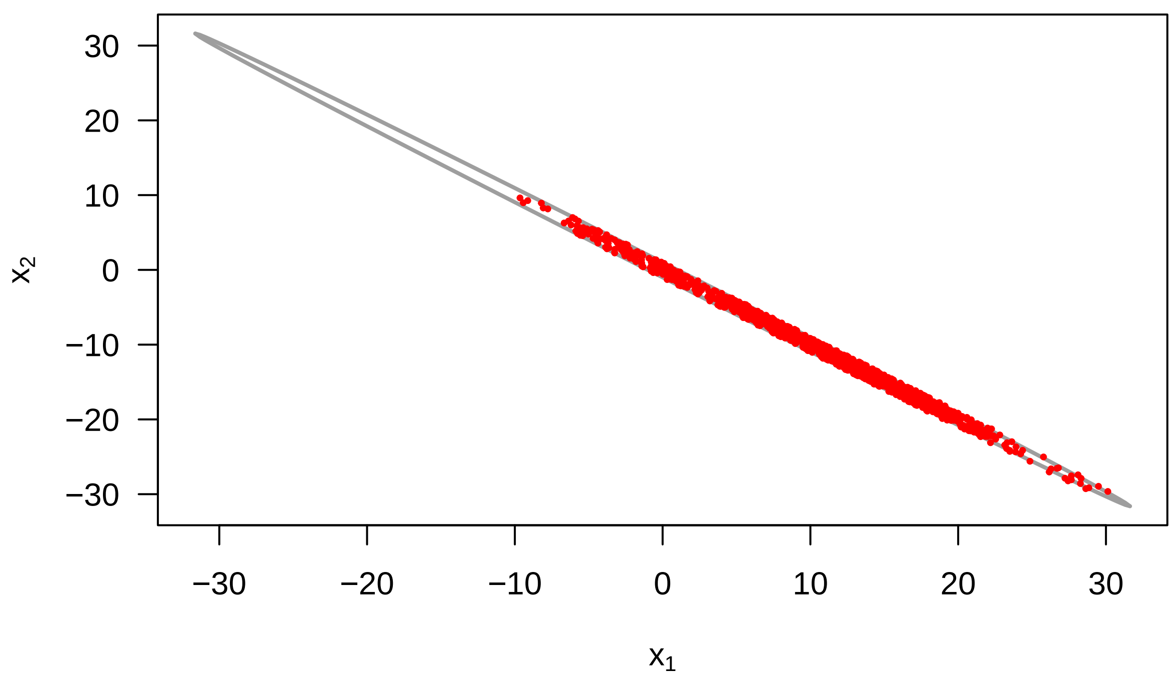

Determine P $(X1 > 0)$ based on the sample and repeat this a few times (you need not to plot the repetitions). What should be the true result for this probability?

$P(X1 > 0) = 0.87$. The true result should be 0.5. Using the same code, we can repeat the experiment a few times and it seems that the result is always close to 0.5, but this is very dependent on the starting values of X1 and X2.

```r
#Determine P(X1 > 0)

means <- c()
for(i in 1:100){
  set.seed(i)
  n <- 1000
  x1 <- 0
  x2 <- 0
  X1 <- c()
  X2 <- c()
  for(i in 1:n){
    x1 <- f1_cond(x2)
    x2 <- f2_cond(x1)
    X1 <- c(X1, x1)
    X2 <- c(X2, x2)
  }
  means <- c(means, mean(X1 > 0))
}
#mean(means)
```

After running the code 100 times, we get a mean of 0.47483.

d.

Discuss, why the Gibbs sampling for this situation seems to be less successful for w = 1.999 compared to the case w = 1.8 from the lecture.

The Gibbs sampling for this situation seems to be less successful for w = 1.999 compared to the case w = 1.8 from the lecture because the boundaries of the region where X has a uniform distribution are very close to each other. This makes it hard for the simulation to cover the whole region because the probability of the simulation to jump from one side of the region to the other is very low. This is why the simulation is very dependent on the starting values of X1 and X2.

e.

We might transform the variable X and generate $U = (U_1, U_2) = (X_1 - X_2, X_1 + X_2)$ instead. In this case, the density of the transformed variable $U = (U_1, U_2)$ is again a uniform distribution on a transformed region (no proof necessary for this claim). Determine the boundaries of the transformed region where U has a uniform distribution on. You can use that the transformation corresponds to $X_1 = (U_2 + U_1)/2$ and $X_2 = (U_2 - U_1)/2$ and set this into the boundaries in terms of X_i . Plot the boundaries for (U_1, U_2). Generate n = 1000 random vectors with Gibbs sampling for U and plot them. Determine $P(X_1 > 0) = P((U_2 + U_1)/2 > 0)$. Compare the results with Part c.

Using the transformation, we get the following expression for $f(u_1, u_2)$:

$f(u_1, u_2) = 1\{u_2^2(0.5 + 0.25 * w) + u_1^2(0.5 - 0.25 * w) < 1\}$

we can then get the conditional by solving the equation $u_2^2(0.5 + 0.25 * w) + u_1^2(0.5 - 0.25 * w) - 1 = 0$ for u_1 and u_2 respectively.

We then obtain that the conditional distribution for u_2 given u_1 is a uniform distribution on the interval $(-\sqrt{\frac{1-u_1^2(0.5-0.25*w)}{0.5+0.25*w}} , \sqrt{\frac{1-u_1^2(0.5-0.25*w)}{0.5+0.25*w}})$ and the conditional distribution for u_1 given u_2 is a uniform distribution on the interval $(-\sqrt{\frac{1-u_2^2(0.5+0.25*w)}{0.5-0.25*w}} , \sqrt{\frac{1-u_2^2(0.5+0.25*w)}{0.5-0.25*w}})$ and

```r
#Gibbs sampling
xv <- seq(-sqrt(4000), sqrt(4000) , by=0.01)   # a range of x1-values, where the term below the root is
plot(xv, xv, type="n", xlab=expression(x[1]), ylab=expression(x[2]), las=1, ylim = c(-3,3))
# ellipse
lines(xv, -sqrt(1-xv^2*0.00025)/(0.99975), lwd=2, col=8)
lines(xv, sqrt(1-xv^2*0.00025)/(0.99975), lwd=2, col=8)

u1_given_ <- function(u2){
  return(runif( 1, min = -sqrt((1-u2^2*0.00025)/0.99975), max = sqrt((1-u2^2*0.00025)/0.99975) ))
}
u2_given_ <- function(u1){
  return(runif(1, min = -sqrt((1-u1^2*0.99975)/0.00025), max = sqrt((1-u1^2*0.99975)/0.00025) ))
}

means <- c()
for(i in 1:100){
  set.seed(i)
  n <- 1000
  u1 <- 0
  u2 <- 0
  U1 <- c()
  U2 <- c()
  for(i in 1:n){
    u1 <- u1_given_(u2)
    u2 <- u2_given_(u1)
    U1 <- c(U1, u1)
    U2 <- c(U2, u2)
  }
  means <- c(means, mean(U1 > 0))
}
mean(means)
```

```
## [1] 0.49937
```

```r
points(X2, X1, col = "red", pch = 20, cex = 0.5)
```

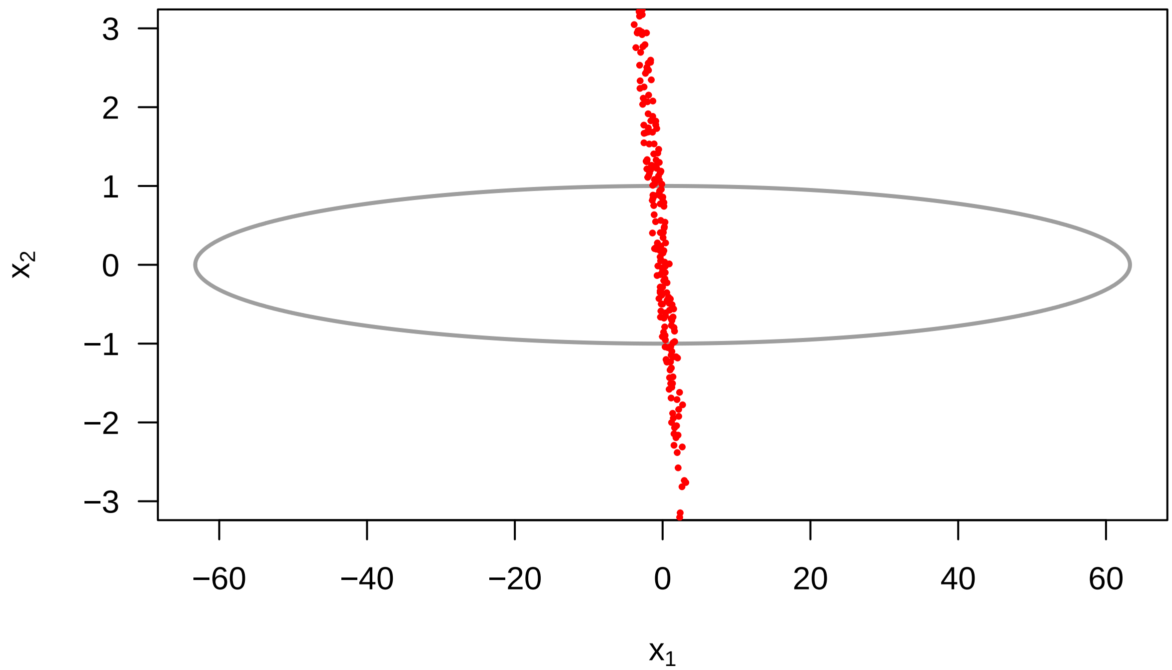
3
2
1
0
−1
−2
−3
x2
−60
−40
−20
0
20
40
60
x1